



BACHELOR'S THESIS IN COMPUTER SCIENCE AND INDUSTRIAL ECONOMICS

UNDERGRADUATE LEVEL 15 CREDITS

A Comparative Evaluation of Open-Source Digital Asset Management Systems

Exploring Organizational and Marketing Criteria for Process and
Marketing Innovation in SMEs

ELLA KARLSSON

School of Industrial Engineering and Management
Royal Institute of Technology (KTH)

Abstract

(?)

• What is the topic area? (optional) Introduces the subject area for the project. • Short problem statement • Why was this problem worth a Master's thesis project? (i.e., why is the problem both significant and of a suitable degree of difficulty for a Master's thesis project? Why has no one else solved it yet?) • How did you solve the problem? What was your method/insight? • Results/Conclusions/Consequences/Impact: What are your key results/conclusions? What will others do based upon your results? What can be done now that you have finished - that could not be done before your thesis project was completed?

Keywords:

Digital Asset Management (DAM), Version Control, Metadata Management, Access Control, SMEs, Workflow Optimization

Sammanfattning

Nyckelord:

Acknowledgments

I would like to thank xxxx for having yyyy.

Contents

Abstract	1
Sammanfattning	1
Acknowledgments	2
List of Figures	5
List of Tables	6
List of Acronyms and Abbreviations	7
1 Introduction	8
1.1 Background	8
1.2 Problem	8
1.3 Purpose	8
1.3.1 Technical Research questions	8
1.3.2 Business Research questions	8
1.3.3 Societal Impact	8
1.3.4 Ethical considerations	9
1.3.5 Sustainability, and social considerations	9
1.4 Goals	9
1.5 Research Methodology	9
1.5.1 Design Science Approach	10
1.5.2 Quantitative and Qualitative Methods	10
1.6 Delimitations	10
1.7 Structure of the thesis	10
2 Background	11
2.1 Artificial Intelligence	11
2.1.1 Object Detection	11
2.1.2 YOLOv11 model	11
2.1.3 Anchor-free detection models	12
2.1.4 The Architecture of a Convolutional Neural Network	13
2.2 Object Detection with YOLOv11	15
2.3 LOSS	16
2.4 Digital Asset Management	16
2.4.1 Choosing a DAM and the key tasks	17
2.4.2 Technological Tools Demand Continuous Organizational Adaptation	17
2.4.3 Why to make our own and not use a service	17
2.4.4 Major background area#1#1	18
2.4.5 What is the YOLO model + how does it work from a higher level perspective	18
2.4.6 The YOLO model	18
2.5 Major background area#2	19
2.5.1 Major background area#2#1	19
2.5.2 Major background area#2#2	19
2.6 Related work	19
2.6.1 Major related work	19
2.6.2 Major related work	19
2.6.3 Minor related work	19
2.7 Summary	19
3 <Engineering-related content, Methodologies and Methods> Use a self-explaining title	20
3.1 Research Process	20
3.2 Research Paradigm	20
3.3 Data Collection	20

3.3.1	Sampling	20
3.3.2	Sample Size	20
3.3.3	Target Population	20
3.4	Experimental design/Planned Measurements	20
3.4.1	Test environment/test bed/model	20
3.4.2	Hardware/Software to be used	20
3.5	Assessing reliability and validity of the data collected	20
3.5.1	Reliability	20
3.6	Validity	20
3.7	Planned Data Analysis	20
3.7.1	Data Analysis Technique	20
3.7.2	Software Tools	20
3.8	Evaluation framework	20
4	[What you did – Choose your own chapter title to describe this]	20
4.1	Hardware/Software design .../ModelSimulation model parameters/...	20
4.2	Implementation .../Modeling/Simulation/...	21
5	Results and Analysis	21
5.1	Major results	21
5.2	Reliability Analysis	21
5.3	Validity Analysis	21
5.4	Discussion	21
6	Conclusions and Future work	21
6.1	Conclusions	21
6.2	Limitations	21
6.3	Future work	21
6.4	Reflections	21
	References	21
	Appendices	24
A	Appendix A: Example Appendix Title	24
B	Appendix B: Another Appendix Example	25

List of Figures

1-1	Sustainable Development Target 9.5 and 12.6	9
2-1	Bounding box for table with legs.	12
2-2	A simplified grayscale sliding kernel operation.	13
2-3	Depicting a Convolutional Neral Netowk	14
2-4	The architecture of YOLOv11, illustrating its three main components: Backbone, Neck, and Head.	14
2-5	Illustrating the five main stages of DAM.	17
2-6	YOLOv11 performance comparison (Ultralytics Inc., 2025).	19
4-1	An example figure in Section.	20
A-1	An example figure in Appendix A.	24

List of Tables

2.1	Summary of YOLO Model Evolution	11
4.1	An example table in Section.	21
5.1	An example table in Section	21
A.1	An example table in Appendix A.	24

List of Acronyms and Abbreviations

AI	Artificial Intelligence
DAM	Digital Asset Management
DSR	Design Science Research
DT	Digital Transformation
ERP	Enterprise Resource Planning
IT	Information Technology
ML	Machine Learning
MCS	Management Control Systems
MDM	Metadata Management
RBAC	Role-based access control
RBV	Resource-Based View
SME	Small and Medium-sized Enterprises
UX	User Experience
VRIN	Valuable, Rare, Inimitable, Non-substitutable
YOLO	You Only Look Once

1 Introduction

To be added

1.1 Background

Digital Asset Management (DAM) emerged in the late 1990s as organizations began grappling with the rapid increase in digital content (Krogh, 2009). Early DAM systems were primarily on-premises solutions designed to store and manage assets such as images, videos, and documents. In the early 2000s, these systems transitioned to cloud-based platforms, offering improved scalability and accessibility (McCain et al., 2021).

More recently, the integration of Artificial Intelligence (AI) and machine learning (ML) has transformed DAM by automating key processes like image tagging, sorting, and categorization. Advanced computer vision techniques now enable systems to analyze and tag images automatically, reducing manual effort and increasing accuracy (Wu et al., 2022).

1.2 Problem

As bespoke manufacturers scale, managing digital assets—spanning product imagery, design renderings, and technical specifications—becomes essential for brand consistency and operational efficiency. However, most DAM solutions, especially open-source systems, lack the necessary automation, posing adoption and maintenance challenges for small and medium-sized enterprises (SMEs) with limited IT infrastructure. Wu et al. studied automated metadata annotation for cultural heritage and found that AI-generated captions often oversimplify context, such as describing a medieval knight merely as a “man on a horse” (Wu et al., 2022). This reflects similar challenges in design-driven manufacturing, where internal product terminology and industry-specific references require more precise and context-aware interpretation.

A core function of DAM is image tagging, sorting, and categorization, directly influencing asset retrievability and structural organization. Although AI has been integrated into some DAM solutions, these implementations typically rely on large pre-trained models that offer broad object classification rather than domain-specific tagging and vocabulary. Recent advancements in computer vision, particularly through algorithms such as YOLO (You Only Look Once), offer an opportunity to overcome these limitations. However, deploying a YOLO-powered system in this domain requires adapting the model to the specific features and vocabulary of the manufacturing sector. Rather than training a model from scratch—a process that demands extensive

annotated data and computational resources—a more feasible approach is to fine-tune a pre-trained model using company-specific data.

1.3 Purpose

The primary aim of this thesis is to assess the feasibility and impact of a YOLO-powered DAM system that has been fine-tuned on company-specific data to address the unique needs of premium manufacturing SMEs. The research will benchmark the performance of this fine-tuned system against a conventional open-source DAM platform (ResourceSpace), focusing on improvements in asset categorization accuracy and retrieval efficiency.

1.3.1 Technical Research questions

- (a) To what extent does fine-tuning YOLOv11 on company-specific data improve metadata accuracy in DAM for manufacturing assets? Can it effectively capture the subtle distinctions of assets?
- (b) What are the trade-offs between the YOLOv11 model and ResourceSpace DAM tagging methods?

1.3.2 Business Research questions

Technological advancements alone do not guarantee successful integration. To complement this, the business perspective assesses the organizational and strategic impact after selecting the preferred DAM system. Specifically:

- (c) How does employee adaptation, the necessity of training, and any role adjustments impact a bespoke manufacturing company?
- (d) (add / incorporate something about process innovation?)
- (e) In what way does improved tagging strengthen brand consistency, customer engagement, and scalability?

1.3.3 Societal Impact

Digital transformation has a significant impact on SMEs. These companies account for approximately 60% of total turnover and value-added contributions in Sweden’s private sector, employing around 65% of the workforce (Tillväxtverket, 2021). The adoption of DAM systems is an integral part of this transformation, improving operational efficiency and reducing manual work, which contributes to broader economic growth. A cost-benefit analysis of 319 SMEs found that digital transformation enhances organizational resilience, reduces operational costs, and improves long-term scalability (Teng et al., 2022).

The stakeholders of this project?

This study is structured around a systematic process encompassing data collection, annotation, model fine-tuning, and testing. These phases represent essential steps that an SME would need to undertake if they were to implement a similar AI-based solution. By addressing both the positive impacts and the possible challenges, the aim is to show if the benefits of adopting this solution justify the necessary investments and efforts. The project's outcomes are expected to contribute to academic knowledge in the field of AI-powered asset management, fostering further innovation.

1.3.4 Ethical considerations

Ethically, the project will investigate issues related to data privacy, transparency, and bias, which are critical in ensuring that automated systems operate fairly and without unintended consequences. These concerns are highlighted in the literature on AI ethics, which emphasizes the need for clear guidelines to mitigate risks associated with autonomous decision-making (Jobin et al., 2019).

1.3.5 Sustainability, and social considerations



Figure 1-1: Sustainable Development Target 9.5 and 12.6

From a sustainability perspective, this research contributes to the United Nations Sustainable Development Goals (SDGs), specifically SDG 9, Industry, Innovation, and Infrastructure, and SDG 12, Responsible Consumption and Production, (United Nations, 2015). In relation to SDG 9, and more precisely target 9.5 as seen in Figure 1-1, the project seeks to enhance scientific research and upgrade the technological capabilities within industrial sectors. Similarly, under SDG 12 target 12.6 also shown in 1-1, this project supports sustainable business practices by optimizing digital asset management. By enhancing asset categorization and retrieval, the system makes it easier for companies to track and store metrics. This dual focus ensures that the technological advancements proposed are not only efficient and innovative but also ethically sound and socially beneficial.

Further reflection will be revisited in Section 6.4.

1.4 Goals

The primary goal is evaluating the feasibility of a YOLO-powered DAM system that has been fine-tuned using company-specific data, in comparison to the open-source solution ResourceSpace. To achieve this, the project has been divided into the following three sub-goals:

1. Dataset Development and Annotation:

Develop a robust methodology for collecting a domain-specific dataset that accurately captures the visual and functional nuances of digital assets in premium manufacturing. The annotation process will involve:

- Using bounding boxes to precisely delineate asset regions.
- Assigning appropriate class labels using a standardized labeling schema to ensure consistency and relevance to the manufacturing domain.

This dataset will serve as the foundation for model fine-tuning.

2. Model Fine-Tuning and Optimization:

Fine-tune a pre-trained YOLO model on the annotated dataset. The objective is to enhance the model's accuracy in tagging, sorting, and categorizing.

- Adjusting hyperparameters and leveraging transfer learning techniques.
- Implementing regularization and validation strategies.

3. Performance Benchmarking and Comparative Analysis:

Benchmark the performance of the fine-tuned YOLO-based DAM system against a conventional open-source DAM called ResourceSpace. Evaluation metrics will include:

- Asset categorization accuracy.
- Retrieval efficiency.
- Overall system usability.

A comparative analysis will be conducted to assess whether the customized system offers significant improvements over traditional solutions. Resulting in practical recommendations and guidelines for manufacturing SMEs considering the adoption of AI-powered DAM.

1.5 Research Methodology

This research employs a mixed-methods approach to address both the technical performance of the system and stakeholder perspectives. Mixed-methods research combines quantitative techniques (e.g., controlled experiments and statistical analyses) with

qualitative techniques (e.g., semi-structured interviews and thematic analysis) to provide a comprehensive evaluation of complex systems (Johnson and Onwuegbuzie, 2004).

Alternative methodologies—such as exclusively quantitative performance evaluations or purely qualitative case studies—were considered but ultimately rejected because they would not fully capture the multifaceted challenges of deploying an AI-powered system in a dynamic industrial environment.

1.5.1 Design Science Approach

Grounded in a pragmatic philosophy that emphasizes practical impact and utility, this study adopts the design science research (DSR) paradigm. DSR is particularly well-suited for technology-driven projects because it promotes the iterative design, development, and rigorous evaluation of IT artifacts to solve real-world problems (Hevner et al., 2004). In this project, the YOLO-powered DAM system represents the artifact developed and refined through iteration.

1.5.2 Quantitative and Qualitative Methods

Controlled experiments will be conducted to measure key performance metrics—such as asset categorization accuracy, retrieval efficiency, and overall system usability. Statistical analysis will be used to validate the improvements brought about by model fine-tuning, following best practices in empirical research (Creswell, 2014; Yin, 2014). Complementing this, qualitative methods will capture contextual insights and stakeholder perspectives. Semi-structured interviews and thematic analysis will be employed to understand user experiences and organizational challenges associated with implementing the DAM system. Moreover, to develop a standardized labeling schema for the dataset, a targeted collaboration with a designated expert from the company will be undertaken. This focused approach is preferred over a large-scale survey. Not all employees interact with digital assets and the expert can ensure domain-specific terminology is accurately captured and applied consistently during annotation.

1.6 Delimitations

This thesis focuses exclusively on evaluating a YOLO-powered digital asset management system for premium manufacturing SMEs. The study is limited to a specific company’s environment and a predefined dataset.

The research investigates only the fine-tuning of an existing pre-trained YOLOv11 model. Training a model from scratch, which requires vast amounts of data and computational resources, is beyond

the scope of this project. Instead of conducting a large-scale survey, the study uses semi-structured interviews with key stakeholders—particularly a designated domain expert—to develop a standardized labeling schema.

This focused approach is chosen because only a few employees directly manage digital assets. The assessment will concentrate on technical performance indicators such as asset categorization accuracy, retrieval efficiency, and overall system usability. Broader issues such as integration with other enterprise systems and macroeconomic impacts are beyond the scope of this project.

1.7 Structure of the thesis

This thesis is organized into the following main chapters, excluding the introductory chapter, references, and appendices; Chapter 2 provides the necessary background and reviews related work, establishing the context for DAM and identifying the key gaps this project addresses. Chapter 3 outlines the methodology—including the design science approach, mixed-methods strategy, data collection, experimental design, and evaluation criteria—used to assess the system. Chapter 4 details the implementation, covering system design, model fine-tuning, dataset development, and the technical setup for testing. Chapter 5 presents the results and analysis, discussing both quantitative metrics and qualitative insights to evaluate whether the project’s goals have been met. Finally, Chapter 6 summarizes the key findings, reflects on the limitations of the study, and outlines potential directions for future work.

2 Background

2.1 Artificial Intelligence

Artificial Intelligence (AI) is a field of computer science that focuses on systems built on algorithms, which are formalized sets of instructions that process input data to produce outputs (Khanam et al., 2024a). Machine Learning (ML), a subset of AI, represents a shift away from manually encoded rules toward data-driven learning. Instead of being explicitly programmed for specific tasks, ML models identify patterns in large datasets and use statistical techniques to make predictions or classify new data.

Khanam et al. (2024a) describe deep learning (DL) as a machine learning approach that utilizes multi-layered computational models to extract patterns from data at varying levels of abstraction. Inspired by the human brain, DL models excel at recognizing intricate patterns in large datasets, making them essential in fields such as image recognition, natural language processing, and autonomous systems (Soori et al., 2023).

Within DL, different neural network architectures are designed to process specific types of data and perform specialized tasks. One of the most effective architectures for structured, grid-like data—such as images and time-series signals—is the Convolutional Neural Network (CNN). CNNs employ convolutional operations to automatically learn spatial hierarchies of features, allowing them to capture patterns and structures in data with high accuracy. As a result, CNNs have become a cornerstone of computer vision, powering applications in object detection, image classification, and other visual recognition tasks (Goodfellow et al., 2016, pp. 326-328).

2.1.1 Object Detection

Object detection involves both the ability to recognize the classes of multiple objects in an image and determining their positions, whereas image classification assigns a single class to the entire image without distinguishing individual objects.

Zhang et al. (2025) outline how DL-based object detection methods are primarily divided into two categories: two-stage and single-stage networks. Two-stage networks, such as Region-Based Convolutional Neural Networks (R-CNNs), rely on generating region proposals before classifying and refining object locations. In contrast, single-stage networks, such as You Only Look Once (YOLO), eliminate this intermediate step by predicting object classes and bounding boxes in a single pass. This approach significantly improves detection speed and efficiency. As Zhang et al. (2025) emphasize, single-stage models have become widely adopted in various industries due to their ability to perform real-time object detection accurately.

2.1.2 YOLOv11 model

The YOLOv11 model, developed by Ultralytics marks the latest milestone in the continuous evolution of the YOLO series, building on a decade of refinement and optimization, as summarized in Table 2.1. Since its introduction by Redmon et al. (2016), it has revolutionized real-time object detection with its single-stage pipeline, offering a faster and more efficient alternative to traditional region-based approaches like R-CNNs.

Release	Key capabilities
V1 JUN 2015	Darknet. A single-stage object detector with basic classification (Redmon et al., 2016).
V2 DEC 2016	Darknet. Object detection. Darknet-19 architecture, anchor boxes, and higher resolution inputs (Redmon and Farhadi, 2016).
V3 MAR 2018	Darknet. Object detection. Darknet-53 network & multi-scale predictions for varying object sizes. (Redmon and Farhadi, 2018).
V4 APR 2020	Darknet. Object detection. Basic object tracking with BCSPDarknet53 and SPP. (Bochkovskiy et al., 2020).
V5 JUN 2020	PyTorch. Object detection. Basic instance segmentation. Multi-GPU support, and exports (Ultralytics, 2020).
V6 SEP 2022	PyTorch. Object detection, instance segmentation, a reparameterizable backbone, anchor aided training (AAT). (Li et al., 2022).
V7 JUL 2022	PyTorch. Object detection, tracking & instance segmentation. (Wang et al., 2022).
V8 JAN 2023	PyTorch. Anchor-free object detection, instance & panoptic segmentation, NVIDIA GPUs, Jetson. (Ultralytics, 2023).
V9 FEB 2024	PyTorch. Anchor-free detection & instance segmentation. PGI for better gradient reliability. GELAN network (Wang et al., 2024b).
V10 MAY 2024	PyTorch. Anchor-free detection & NMS-free training (Wang et al., 2024a).
V11 SEP 2024	PyTorch. Anchor-free & oriented object detection (OBB), instance segmentation, pose estimation. (Ultralytics Inc., 2025).
V12 FEB 2025	PyTorch. Anchor-free detection, OBB, instance segmentation, Area Attention Mechanism, pose estimation, R-ELAN. (Ultralytics Inc., 2025).

Table 2.1: Summary of YOLO Model Evolution

Early versions of YOLO were built on the Darknet framework, developed by Joseph Redmon, with core implementations written in C and CUDA for fast GPU execution. A framework is a pre-built structure that simplifies software development by providing reusable code, tools, and libraries allowing

developers to focus on higher-level abstraction. As shown in Table 2.1, the transition to PyTorch occurred with YOLOv5, developed by Ultralytics. PyTorch, originally introduced by Facebook AI Research (FAIR), offered a more flexible and scalable environment, facilitating development in Python and enhancing integration with mainstream deep learning research (Ultralytics, 2020).

Sapkota et al. (2025) conducted a comprehensive review of YOLO-based object detection applications, highlighting its extensive adoption across multiple domains, including healthcare (e.g., pill identification, diagnostics), surveillance (e.g., face mask detection, home security), autonomous vehicles, and industrial quality control. The study underscores YOLO’s efficiency in real-time processing, making it a preferred choice for applications requiring rapid inference.

While YOLO excels in speed, its grid-based detection approach and anchor-free methodology maintained in YOLOv6 and subsequent models introduce inherent limitations. Both Sapkota et al. (2025) and He et al. (2024) note that, despite its computational efficiency, YOLO may struggle with fine-grained detail detection, making it less suitable for tasks requiring high-resolution texture analysis, such as road damage assessment or material surface inspection (Angulo et al., 2019). While this thesis primarily addresses the application of YOLO within DAM in bespoke manufacturing, insights into the limitations remain highly relevant, particularly in scenarios where accurate detection and classification of subtle material textures effect performance.

The trade-off between speed and accuracy is further emphasized in comparative analyses, such as Rane (2023), which contrasts YOLO with Faster R-CNN. While YOLO excels in inference speed—making it well-suited for real-time applications such as inventory management, checkout automation, and e-commerce visual search—Faster R-CNN offers superior object localization and classification accuracy. This aligns with the findings of Sapkota et al. (2025), making it the preferred choice for scenarios demanding precise differentiation and high recall, such as medical imaging. However, Faster R-CNN’s reliance on a region proposal network (RPN) results in significantly higher computational demands, limiting its viability for real-time deployment (Rane, 2023).

In contrast, the study by Karbouj et al. (2024) on object detection for screw head identification in disassembly systems presents a different perspective. Their findings demonstrate that YOLOv5 outperforms Faster R-CNN across multiple key metrics, including precision, recall, inference speed (FPS), and training efficiency. This discrepancy arises from the nature of the application and dataset size. As

discussed by Rane (2023) Faster R-CNN tends to perform better in tasks requiring high-detail object recognition. The RPN helps it generalize more effectively when training data is limited, making it particularly useful for small datasets with high precision requirements. Conversely, YOLO’s ability to efficiently learn broad patterns makes it a superior choice for large-scale, high-variance datasets. The findings of Karbouj et al. (2024) reinforce this perspective, demonstrating YOLOv5’s balance between computational speed and adaptability, making it particularly effective in real-time, resource-constrained environments.

2.1.3 Anchor-free detection models

A bounding box defines an object’s position and size within an image using four coordinates. In object detection, it is paired with a class label and a confidence score, indicating both the object’s category and the model’s certainty in its prediction. These boxes act as ground-truth references in training data, helping models learn to localize objects accurately. (Li et al., 2022). The prediction represents the final output of an object detection model as illustrated in Figure 2-1.




Figure 2-1: Bounding box for table with legs.

Vina (2024) highlights the shift from anchor-based to anchor-free object detection as a major advancement in the field. Traditional anchor-based detectors, such as YOLOv4 and its predecessors in Table 2.1, rely on predefined anchor boxes—fixed-size reference shapes placed across an image at different aspect ratios—to estimate object locations. The model does not predict bounding boxes directly but instead modifies the closest anchor to better fit detected objects. Anchor-free models simplify detection and improve speed—critical for real-time tasks like autonomous driving and surveillance. Their key-point-based approach enhances flexibility, making them better at detecting small, irregular, or occluded objects, especially in cluttered environments where anchor-based methods struggle (Wang et al., 2024b).

2.1.4 The Architecture of a Convolutional Neural Network

According to (Prince, 2023, p. 163) digital images possess three essential properties that demand specialized model architectures. Images are high-dimensional. For example, a typical one is 224×224 pixels with three color channels (RGB), amounting to 150,528 input dimensions. Fully connected networks would require an enormous number of parameters to process such data. Moreover, nearby pixels tend to be correlated, and the overall content of an image remains recognizable even with small shifts.

At a fundamental level, CNNs process input through sequential stages, using convolution to detect features, pooling to reduce dimensionality, and activation functions to introduce non-linearity. In this setup, convolution identifies spatial features such as textures, lines and color variations in the input. With effective training, the network learns to recognize these attributes regardless of their location within an image (Verdhan, 2021, Chapter 2).




Figure 2-2: A simplified grayscale sliding kernel operation.

As described by (Prince, 2023, p. 170) and illustrated in Figure 2-2, a convolution operation applies a 3×3 kernel K with learned weights to an input I in order to generate a feature map O . At each spatial position, the kernel computes a weighted sum over a local 3×3 patch of the input. Including a bias term b and an activation function $a(\cdot)$, the operation is given by:

$$O_{ij} = a \left(b + \sum_{m=1}^3 \sum_{n=1}^3 I_{i+m-2, j+n-2} \cdot K_{mn} \right), \quad (1)$$

where O_{ij} denotes the output at position (i, j) . The indices $i + m - 2$ and $j + n - 2$ center the kernel over the input. When no padding is applied, the output dimensions are reduced compared to the input. However, if a padding of size p is introduced (e.g., $p = 1$ for a 3×3 kernel), the input is effectively extended so that the spatial dimensions of O can match those of I .

As described by Verdhan (2021), Figure 2-2 illustrates a simplified example of a convolution operation using a 3×3 kernel K , also known as a filter or feature detector, which contains learned

weights. The kernel slides across the 5×5 input matrix I with a stride of $s = 1$. At each step, it overlaps a patch of I —known as the receptive field—and computes the dot product between its weights and the patch. This operation follows Equation 2:

$$O_{ij} = \sum_{m=1}^3 \sum_{n=1}^3 I_{i+m-1, j+n-1} \cdot K_{mn}, \quad (2)$$

where O_{ij} is the element at the i -th row and j -th column of the feature map O , (with $1 \leq i, j \leq 3$). Notice that because no padding is applied, the dimensions of O are reduced compared to the input.

Digital images are composed of pixels that encode brightness and color information at specific spatial locations. As explained by Khanam et al. (2024b), in a standard 8-bit image each pixel is represented by an intensity value ranging from 0 (darkest) to 255 (brightest).

Figure 2-2 demonstrates a convolution operation using a simplified 5×5 single-channel (grayscale) matrix. Most images are represented as three-dimensional tensors with dimensions $H \times W \times C$, where H is the height, W is the width, and C is the number of channels. For example, a standard RGB image has $C = 3$ channels corresponding to red, green, and blue (Prince, 2023, p. 170). One can visualize this as three separate layers—one for each color—stacked along the depth dimension, much like stacking three copies of I inside Figure 2-2 on top of each other.

Similarly, each convolution kernel is represented as a three-dimensional tensor of size $K \times K \times C$. The depth of the kernel must match the depth of the input (C) to perform element-wise multiplication correctly, often with an added bias to produce a single scalar output. Verdhan (2021)

Early layers of a CNN learn low-level features like edges and curves, while deeper layers combine these to form high-level representations such as faces or objects. However, as the network deepens, the number of dimensions and overall complexity increase, especially when image augmentations (like rotations) alter the feature maps. To manage this complexity, downsampling is used, which reduces the resolution of the feature maps. This is achieved through a pooling layer placed after the convolutional layers. Pooling layers operate on each feature map individually, reducing their dimensions—often halving them—by summarizing regions. Two common pooling methods are average pooling, which computes the mean value, and max pooling, which selects the maximum value from each region (Verdhan, 2021).

These layers apply small, learnable filters, also known as kernels, which systematically scan the image, detecting patterns such as edges, textures, and

shapes. Each kernel responds to specific features in the input, enabling the network to learn hierarchical representations through multiple layers (Alif and Hussain, 2025).

In digital representation, an image is stored as a numerical matrix of pixel values. A color image is typically encoded as a three-dimensional tensor ($H \times W \times C$), where C represents the number of color channels—commonly three for RGB images.

Unlike traditional fully connected networks, CNNs utilise spatially local connections and weight sharing to improve computational efficiency and reduce the number of parameters

As described by

When processing images CNNs use convolutional layers. Scanning the image using filters also called kernels that can detect features. Each kernel enables the model to learn patterns by activating a specific pattern in the input image.

Following convolution activation functions such as ReLU are used. It outputs only positive numbers, zero otherwise. This helps the model learn nonlinear data for pattern recognition.

Pooling is a downsampling operation where downsampling means to reduce the dimensionality of the feature map extracted. This reduces the computational cost and also overfitting.

In the end there are fully connected (FC) layers similar to those of a neural network. These interpret the features.

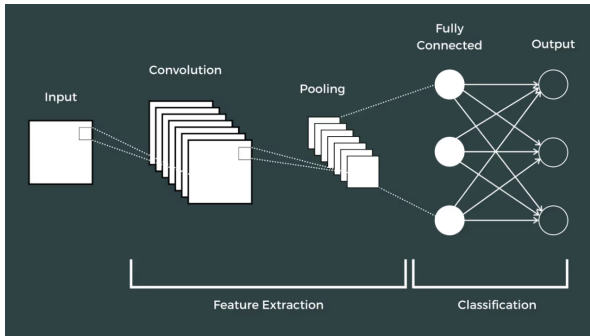


Figure 2-3: Depicting a Convolutional Neural Network

Architecture Both the YOLOv11 and V12 model follows the standard three-part structure of the YOLO family: Backbone, Neck, and Head, as shown in Figure 2-4. According to Hidayatullah et al. (2025), the Backbone extracts features using convolutional layers and downsampling, generating hierarchical feature maps. The Neck refines these features through the SPPF block for multi-scale detection and the C2PSA module to enhance the recognition of small and occluded objects. Upsampling and feature concatenation further improve resolution and information retention. Finally, the Head produces the model's output, predicting class probabilities and

bounding boxes across three detection layers (small, medium and large), each specialized for different object sizes (Hidayatullah et al., 2025).

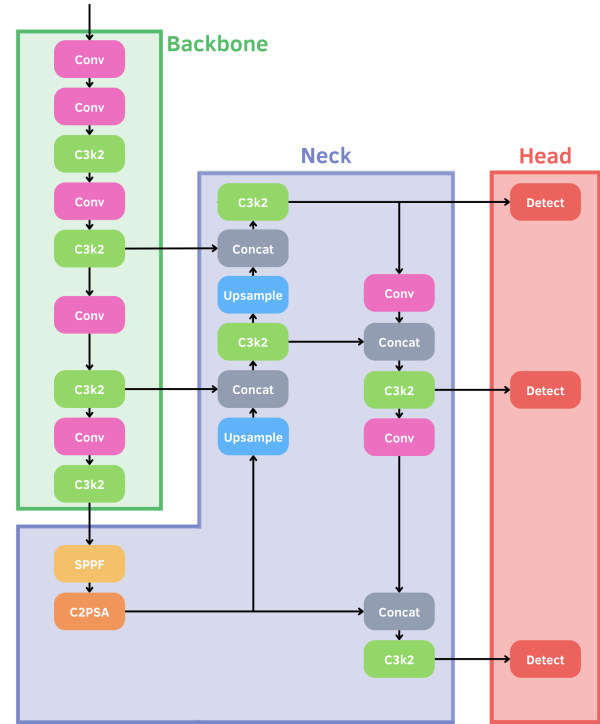


Figure 2-4: The architecture of YOLOv11, illustrating its three main components: Backbone, Neck, and Head.

and the Head Create bounding boxes and pair it to a class.

step 1. Overlay image with a grid in size $s \times s$. Each grid cell produces 2 things: - 1. a set of bounding boxes centered on a point inside the grid with confidence scores that an object exists inside each bounding box

- 2 a class probability map for each cell: which tells what object class is most likely to be in that cell given that an object exists within the cell.

It combines this info to yield the object detections TEST

(Alif and Hussain, 2025)

On one hand, YOLO's real-time inference speed, lightweight architecture, and open-source nature make it a compelling choice for applications requiring fast, on-device processing.

As for relating to this thesis. there is limited research on the use of YOLO directly relating for Digital Asset Management (DAM) applications. with only one identified study—Angulo et al.

YOLOv1 unified architecture divided the image into a grid, predicting bounding boxes and class probabilities directly for each cell, enabling end-to-end learning. Since then it has been applied in healthcare including but not limited to pill

identification, diagnostic processes and treatment outcomes, surveillance for social distancing and face mask detection during pandemics and identification, and autonomous vehicle checks and also in detect and classify crops Sapkota et al. (2025).

There is also use of YOLO in industry. inspection processes to detect defects and anomalies and moreover ensuring quality control in manufacturing and production.

There is limited sources of
citeSapkota2025YOLOv11.

Since then a lot has happened. In the beginning. YOLOv1 - v4 relied on anchor boxes.

The improvements of YOLOv11 YOLOv11 outperformed previous versions in mean average precision (mAP), recall, and precision, demonstrating superior object detection performance. The recall rate, which measures how well the model detects all ground-truth objects, was highest for YOLOv11 (64.8%). YOLOv11 also exhibited fewer false detections compared to its predecessors. YOLOv11 displayed higher attention concentration on relevant objects, meaning it focused better on wires and transformers, reducing errors in object localization.

Introduction of the C3K2 module (a new CSP bottleneck layer with smaller convolutional layers), improving feature extraction and detection efficiency. Integration of the C2PSA module, which enhances spatial and channel information processing for more accurate detections. Optimized adaptive anchor box mechanism, improving the model's adaptability across different object sizes and datasets.

The kernel operation At the core of CNNs is the convolution operation, a specialized linear process that replaces the typical matrix multiplication found in other neural network layers. Convolution involves computing a weighted average over the input, where a kernel (or filter) systematically processes the input data to produce a feature map. The figure 2-2 showcases the convolutional operation.

Convolutional filters, also called Local Receptive Fields or kernels are used to detect patterns or features in input data.

The pixel values of an image are the features, because each pixel value represents the intensity of light at a given location. We want to learn the relationship between those features in an image in order to detect objects, faces etc.

For example, in image processing, filters might be designed to detect edges, corners, or textures. In deep learning, the weights of these filters are learned automatically through training on large datasets.

2.2 Object Detection with YOLOv11

state-of-the-art one-stage object detection algorithm renowned for its efficiency and simplicity

construction of a object detection dataset

image preprocessing,

model training using the object detection training dataset,

and validation of results using a verification dataset

YOLO's backbone network has undergone substantial advancements, integrating deeper feature fusion and multiscale feature extraction to enhance its capability for power equipment object detection.

Starting from YOLOv8 [13], the series adopted an anchor-free mechanism for the first time, allowing greater adaptability to detect power equipment targets of varying sizes.

Since YOLOv5 [12], the algorithm has significantly improved detection efficiency and accuracy through the introduction of the CSPNet framework, which optimizes feature propagation and network capacity

updates to the YOLO series have included innovative enhancements to the loss function, further refining the model's detection precision. While the original YOLO algorithm offered remarkable detection speed, its accuracy lagged behind two-stage detection algorithms. H

The incorporation of a spatial pyramid pooling (SPP) layer into the backbone network further expanded the model's receptive field, enhancing its feature extraction capabilities. YOLOv5 advanced these capabilities by adopting the C3 module in its backbone network, which reduced computational complexity and improved inference speed. It also introduced Mosaic data augmentation, particularly Mosaic4, which combines and transforms four images randomly to enhance feature representation and model learning. Adaptive anchor box optimization was added, enabling the model to better handle objects of different sizes. YOLOv8 refined the architecture further by replacing the C3 module with the C2f module, enhancing feature extraction efficiency.

It also introduced an Anchor-Free detection mechanism to improve the detection of small targets. The Mosaic augmentation process was optimized to exclude its use in the final ten training epochs, thereby improving model generalization. Additionally, task-specific loss optimizations were integrated to further enhance detection performance. YOLOv9 [16] introduced progressive gradient integration (PGI), addressing limitations of deep supervision in extremely deep architectures and making lightweight architectures more practical. A new network architecture, called generalized high-efficiency layer aggregation network (GELAN), was proposed. GELAN integrates cross stage partial network (CSPNet) and efficient layer aggregation network (ELAN) designs,

balancing model lightweight design, inference speed, and accuracy. Crossstage partial connections were employed to link feature maps across stages, enriching semantic information and improving

The most recent iteration, YOLOv11, replaced the C2f module with the C3K2 module, a custom CSP bottleneck layer featuring two smaller convolutional layers, improving processing speed without compromising performance. While retaining the SPPF module from YOLOv8, YOLOv11 introduced the C2PSA module, which integrates channel and spatial information with multi-head attention mechanisms for more efficient feature extraction. An adaptive anchor box mechanism was also refined to optimize configurations across diverse datasets, boosting detection accuracy. Beyond object detection in power equipment, YOLOv11 extends its capabilities to instance segmentation, image classification, pose estimation, and oriented bounding box detection (OBB), add

Among these, You only look once (YOLO) , a real-time object detection algorithm, has gained widespread attention. Unlike traditional methods, YOLO eliminates the need for pre-generated candidate regions, directly predicting the class and location of targets within an image. Since its inception in 2015, YOLO has undergone significant advancements, with the latest version, YOLOv11, demonstrating substantial improvements in detection speed and performance

Architectures within the object detection domain can be classified into single-stage or two-stage detectors

YOLO significantly enhances the speed, efficiency, and accuracy of medical object detection compared to traditional methods.

Typical neural network: Input neuron (each connected to each next layer)- hidden layer - output layers

In a convolutional network it is not mandatory all neurons are connected to each in the next hidden layer.

Filters: the fixed square called a patch or local receptive field

Feature map: The feature map is the output of one filter applied to the previous layer.

The filter moves across the input layer. (multiply the values within the filter with the values in the input layer). A new matrix with less dimensions is compartmentalized

input layer is called local receptive fields.

Activation and Pooling layers: Activation: transforming to the output using Activation functional like Resulting(discard the negative values and replace them with zeros)

Pooling: The feature map dimensionality is reduced using pooling (only important features remain... max, min pooling etc i.e to the only

largest)

1. convolutional layers 2. pooling layers 3. fully connected layers

2.3 LOSS

The YOLOv11 object detection method enhances its performance by minimizing a comprehensive loss function that integrates multiple components. This loss function encompasses distributed focal loss, bounding box regression loss, and class probability loss. The optimization process involves combining these individual loss components and employing advanced optimization algorithms to refine the model's performance in object detection tasks

architecture: serves as the foundation for extracting multi-scale feature maps from input images. This is achieved through a series of convolutional layers and specialized modules, designed to generate feature maps at varying resolutions. These feature maps capture the spatial and semantic information necessary for subsequent processing. The Neck functions as an intermediate stage, tasked with aggregating and enhancing features from multiple scales before passing them to the Head network for prediction. This process often involves upsampling and concatenation of feature maps, enabling the model to efficiently capture and utilize multi-scale information. The Neck plays a crucial role in bridging the Backbone and Head components, enhancing feature expressiveness and supporting robust predictions. The Head is responsible for generating the final outputs, including object bounding boxes and object classes labels. It processes the enriched feature maps from the Neck to predict object locations and classifications with high precision. In summary, the Backbone extracts essential feature representations, the Neck aggregates and refines these features across scales, and the Head produces the final predictions. The Neck serves as a pivotal link, combining multi-scale features from the Backbone and augmenting their expressiveness through upsampling and concatenation, thereby providing a strong foundation for the Head's accurate and reliable predictions.

2.4 Digital Asset Management

Krogh (2009) describes DAM as an essential framework for protecting, organizing, and prolonging the usability of digital files by emphasizing metadata, suitable file formats, and efficient workflows. As shown in Figure 2-5, five interconnected stages—creation, management, distribution, archiving, and retrieval—collectively ensure that digital assets remain discoverable and relevant long after their initial production.

Although Krogh does not explicitly align his approach with the Resource-Based View (RBV), his

emphasis on preserving assets as integral organizational resources parallels RBV's tenet that competitive advantage relies on valuable, rare, inimitable, and non-substitutable (VRIN) capabilities (Barney, 1991). By structuring DAM processes around rigorous metadata management, secure storage, and ongoing accessibility, organizations can treat their digital repositories as strategic assets, safeguarding long-term benefits that are difficult for competitors to replicate.

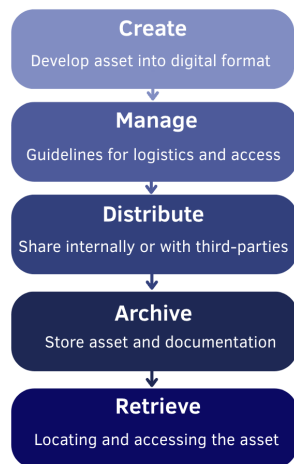


Figure 2-5: Illustrating the five main stages of DAM.

2.4.1 Choosing a DAM and the key tasks

What tools are available in DAM? Benchmark? What are the most important shit in it? What do most companies need? What do they usually have and how or why do they choose to adopt a DAM

A missing perspective is

2.4.2 Technological Tools Demand Continuous Organizational Adaptation

Love and Matthews (2019) identify a critical gap in the construction industry: knowing “why” to adopt digital technologies is relatively straightforward, but knowing “how” to translate technological potential into real value remains largely underexplored. Their case studies underscore the fact that digital transformation does not happen automatically; organizations must actively invest in processes such as benefits management and the development of a Business Dependency Network (BDN) to realize tangible gains from their digital initiatives (Love and Matthews, 2019).

In a broader context, Hanelt et al. (2020) posit that digital transformation (DT) goes beyond any single disruptive episode; it is a continual, structural adjustment propelled by digital technologies. Their systematic review of 279 peer-reviewed articles frames DT across three dimensions—Contextual

Conditions (e.g., technological advances, shifting consumer habits), Mechanisms (e.g., the innovative strategies organizations adopt), and Outcomes (e.g., changes to organizational structures and industry norms). By proposing a typology that spans technology impact, compartmentalized adaptation, systemic shift, and holistic co-evolution, they challenge the idea of one-off change, advocating instead for an iterative, agile approach to transformation (Hanelt et al., 2020).

Taken together, these two perspectives highlight that while there is strong motivation to deploy new technologies (“why”), sustained, organization-wide benefits only materialize when there is a concerted effort to integrate, evaluate, and adapt these digital tools in an ongoing manner (“how”). Both studies imply that true success hinges on long-term structural and cultural shifts rather than static, one-off solutions.

that the promise of DAM is not unlocked simply by adopting new technology but only when companies embrace two fundamental principles. First, that technology alone does not create value but must be accompanied by organizational process reengineering, and second, that the benefits of DAM are maximized only through continuous strategic governance to monitor and sustain its impact

A missing perspective in

Nevertheless, some scholars argue that resource possession alone does not guarantee successful digital transformation. Civelek et al. (2023) found no significant link between dynamic capabilities—a key aspect of RBV that involves adapting, integrating, and reconfiguring resources—and successful digital transformation among Czech manufacturing SMEs. Their findings suggest that merely possessing dynamic capabilities is insufficient for digital transformation unless supported by complementary factors such as digital literacy and IT infrastructure maturity.

2.4.3 Why to make our own and not use a service

Bynder

Adobe Experience Manager

Cloudinary: custom pricing for enterprise solutions.

Adobe Sensei essentially means auto-tagging images based on recognizable generic objects, scenes, and concepts. It typically uses generalized, pre-trained models that identify common objects’

most DAM platforms rely on third-party integrations for company-specific tagging

Clarifai Custom Models Provides APIs that integrate into DAM platforms.

Amazon Rekognition Custom Labels: Pay-per-use
Google Vertex AI (formerly AI Platform Vision)

Pricing depends on training hours and predictions
Custom vision API: Trained specifically on your images and product labels.

Microsoft Azure Custom Vision: Training: 20 dollar per compute hour

Integrates via REST API to enhance tagging accuracy in DAM solutions.

CV consulting

Image annotation

Different types of CV:

?? is an image ?? is a table

2.4.4 Major background area#1#1

Recent studies have demonstrated the effectiveness of various AI techniques in image tagging. Zhang et al. (2019) showcased the application of convolutional neural networks (CNNs) for automatic image classification in DAM systems, achieving an accuracy of 92% on a diverse dataset of digital assets

This work was further extended by Li and Chen (2020), who integrated attention mechanisms into CNNs, improving the model's ability to focus on salient features and increasing tagging accuracy to 95%

The YOLO (You Only Look Once) algorithm has also been applied successfully in DAM contexts. Wang et al. (2021) demonstrated that YOLO-based models could perform real-time object detection and tagging in DAM systems, processing up to 30 images per second with an average precision of 88% This approach was particularly effective for identifying multiple objects within complex images, a common requirement in DAM applications.

Transformer-based models have recently gained traction in image tagging for DAM systems. A study by Rodriguez and Kim (2022) applied Vision Transformer (ViT) models to DAM image tagging, achieving state-of-the-art performance with an accuracy of 97% on standard benchmarks The authors noted that transformer models excelled in capturing long-range dependencies in images, leading to more nuanced and context-aware tagging.

While AI-powered image tagging offers significant benefits, it also presents several challenges. Data requirements pose a significant hurdle, as highlighted by Brown et al. (2020), who found that AI models required at least 10,000 labeled images per category for optimal performance in domain-specific DAM applications

Error rates and handling domain-specific content remain ongoing challenges. A comprehensive study by Thompson et al. (2021) analyzed error patterns in AI-powered image tagging across various industries, revealing that error rates increased significantly (up to 25%) when dealing with highly specialized or technical imagery

To address this issue, Nguyen and Patel (2022) proposed a hybrid approach combining pre-trained models with domain-specific fine-tuning, reducing

error rates by 40% in niche industries such as medical imaging and aerospace engineering

Despite these challenges, the benefits of AI-powered image tagging in DAM systems are substantial. A large-scale study by Garcia et al. (2023) across 500 organizations found that implementing AI-powered tagging led to a 60% reduction in manual tagging time and a 35% improvement in asset discoverability

Entangled states are an important part of quantum cryptography, but also relevant in other domains. This concept might be relevant for neutrinos, see for example [2].

2.4.5 What is the YOLO model + how does it work from a higher level perspective

Object detection algorithm. It locates object in an image It is in two stage detection, it is faster than two stage. It is just an algorithm... so you have to do a lot around it?

Scheme

Architecture

Loss Function

2.4.6 The YOLO model

As demonstrated in table 2.1 the YOLO series has evolved significantly since its inception, introducing progressive improvements in object detection, computational efficiency, and feature extraction. YOLOv11 is the best choice for the project due to its superior accuracy, efficiency, and versatility. As Khanam and Hussain (2024) highlight, its architectural upgrades enhance feature extraction while minimizing computational costs, making it ideal for real-time applications requiring both speed and precision (Khanam and Hussain, 2024).

Beyond object detection, YOLOv11 supports instance segmentation, pose estimation, and oriented object detection, offering greater adaptability to the project's needs. Its optimized balance of accuracy and processing speed ensures strong performance across different computing environments, from edge devices to high-performance systems, making it the most effective solution

The selection of YOLOv11 for the project is driven by its superior architectural enhancements, versatile task support, and optimized balance between accuracy and efficiency. Each version has incorporated refinements aimed at enhancing real-time performance, with YOLOv11 representing the most advanced iteration to date (Khanam and Hussain, 2024).

Computational methods are increasingly used as a




Figure 2-6: YOLOv11 performance comparison (Ultralytics Inc., 2025).

third method of carrying out scientific investigations. For example, computational experiments were used to find the amount of wear in a polyethylene liner of a hip prosthesis in [3].

2.5 Major background area#2

The application of AI-powered image tagging in DAM systems extends beyond large corporations to small and medium-sized enterprises (SMEs), particularly in premium manufacturing sectors. A case study by Hoffmann and Schulz (2022) examined the implementation of AI-powered DAM in a high-end carpentry company similar to Veermakers. The study found that AI-assisted tagging improved product catalog management efficiency by 45% and reduced time-to-market for new designs by 30%.

However, Chen et al. (2023) noted that SMEs in specialized manufacturing often face unique challenges in adopting AI-powered DAM systems, including limited datasets and highly specific visual content. To address these issues, the authors proposed a transfer learning approach, adapting pre-trained models to domain-specific tasks with minimal additional data, achieving a 75% reduction in required training data while maintaining 90% of the original accuracy.

While academic research has made significant strides in advancing AI-powered image tagging techniques, commercial implementations often lag behind in adopting cutting-edge methods. A comprehensive survey by Martinez and Lee (2022) of 50 leading DAM vendors revealed that only 30% had implemented transformer-based models, despite their superior performance in academic studies. The authors attributed this gap to factors such as implementation complexity, computational requirements, and the need for backward compatibility with existing systems.

2.5.1 Major background area#2#1

The integration of AI-powered image tagging in DAM systems raises important ethical, societal, and legal considerations. Privacy concerns are paramount, as highlighted by a study by Johnson and Smith (2022), which found that 35% of automatically generated tags in a sample of 10,000 images

contained potentially sensitive information²². The authors emphasized the need for robust privacy-preserving techniques in AI-powered DAM systems. Algorithmic bias presents another significant challenge. Research by Park et al. (2023) revealed systematic biases in AI-generated tags across gender, ethnicity, and age dimensions, with error rates up to 20% higher for underrepresented groups. This study underscores the importance of diverse and representative training data in mitigating bias in AI-powered DAM systems.

2.5.2 Major background area#2#2

The potential impact on employment is also a concern. While Garcia et al. (2023) found that AI-powered tagging led to significant efficiency gains, they also noted a 15% reduction in human tagging roles across surveyed organizations. However, the same study observed a 10% increase in higher-skilled positions related to AI model management and quality assurance, suggesting a shift rather than a net loss in employment.

2.6 Related work

2.6.1 Major related work

Do not use the title of the paper/book/... as the title of the section. Instead summarize what the contribution of this work is in your own words.

Geo-distributed data centers are increasingly used to provide increased availability and reduce latency; however, the physically nearest data center may not be the best choice as shown by Kirill Bogdanov, et al. in their paper “The Nearest Replica Can Be Farther Than You Think” [4]. Exploring decentralized approaches to AI model training, allowing organizations to collaborate on improving tagging accuracy while preserving data privacy.

2.6.2 Major related work

Carrier clouds have been suggested as a way to reduce the delay between the users and the cloud server that is providing them with content. However, there is a question of how to find the available resources in such a carrier cloud. One approach has been to disseminate resource information using an extension to OSPF-TE, see Roozbeh, Sefidcon, and Maguire [5].

2.6.3 Minor related work

Do not use the title of the paper/book/... as the title of the section. Instead summarize what the contribution of this work is in your own words.

2.7 Summary

It is nice to bring this chapter to a close with a summary. For example, you might include a table that summarizes the ideas of others and the advantages and disadvantages of each – so that later you can compare your solution to each of these. This

will also help guide you in defining the metrics that you will use for your evaluation.

3 <Engineering-related content, Methodologies and Methods> Use a self-explaining title

The contents and structure of this chapter will change with your choice of methodology and methods. For example, if you have implemented an artifact, what did you do and why? How will your evaluate it.

Describe the engineering-related contents (preferably with models) and the research methodology and methods that are used in the degree project. Give a theoretical description of the scientific or engineering methodology are you going to use and why have you chosen this method. What other methods did you consider and why did you reject them. In this chapter, you describe what engineering-related and scientific skills you are going to apply, such as modeling, analyzing, developing, and evaluating engineering-related and scientific content. The choice of these methods should be appropriate for the problem. Additionally, you should be consciousness of aspects relating to society and ethics (if applicable). The choices should also reflect your goals and what you (or someone else) should be able to do as a result of your solution - which could not be done well before you started. The purpose of this chapter is to provide an overview of the research method used in this thesis. Section 3.1 describes the research process. Section 3.2 details the research paradigm. Section 3.3 focuses on the data collection techniques used for this research. Section 3.4 describes the experimental design. Section 3.5 explains the techniques used to evaluate the reliability and validity of the data collected. Section 3.6 describes the method used for the data analysis. Finally, Section 3.7 describes the framework selected to evaluate xxx.

3.1 Research Process

Image of: steps conducted to do the research Fig: research processes

3.2 Research Paradigm

3.3 Data Collection

(This should also show that you are aware of the social and ethical concerns that might be relevant to your data collection method.)

3.3.1 Sampling

1. Aa 2. Bb 3. Cc

3.3.2 Sample Size

3.3.3 Target Population

3.4 Experimental design/Planned Measurements

3.4.1 Test environment/test bed/model

Describe everything that someone else would need to reproduce your test environment/test bed/model/...

3.4.2 Hardware/Software to be used

3.5 Assessing reliability and validity of the data collected

3.5.1 Reliability

How will you know if your results are reliable?

3.6 Validity

How will you know if your results are valid?

3.7 Planned Data Analysis

3.7.1 Data Analysis Technique

3.7.2 Software Tools

3.8 Evaluation framework

4 [What you did – Choose your own chapter title to describe this]

What have you done? How did you do it? What design decisions did you make? How did what you did help you to meet your goals?

4.1 Hardware/Software design .../ModelSimulation model parameters/...

Figure 4-1 shows a simple icon for a home page. The time to access this page when served will be quantified in a series of experiments. The configurations that have been tested in the test bed are listed in Table 4-1.



Figure 4-1: An example figure in Section.

4-1 is an image 4.1 is a table

Column 1	Column 2
Data 1	Data 2
Data 3	Data 4

Table 4.1: An example table in Section.

4.2 Implementation .../Modeling/Simulation/...

5 Results and Analysis

In this chapter, we present the results and discuss them.

Keep in mind: How you are going to evaluate what you have done? What are your metrics? Analysis of your data and proposed solution Does this meet the goals which you had when you started?

5.1 Major results

Some statistics of the delay measurements are shown in Table 5-1. The delay has been computed from the time the GET request is received until the response is sent.

Column 1	Column 2
Data 1	Data 2
Data 3	Data 4

Table 5.1: An example table in Section

5.1 is a table

5.2 Reliability Analysis

LALALA

5.3 Validity Analysis

LALALA

5.4 Discussion

6 Conclusions and Future work

«Add text to introduce the subsections of this chapter.»

6.1 Conclusions

Describe the conclusions (reflect on the whole introduction given in Chapter 1). Discuss the positive effects and the drawbacks. Describe the evaluation of the results of the degree project. Did you meet your goals? What insights have you gained? What suggestions can you give to others working in this

area? If you had it to do again, what would you have done differently?

6.2 Limitations

What did you find that limited your efforts? What are the limitations of your results?

6.3 Future work

Describe valid future work that you or someone else could or should do. Consider: What you have left undone? What are the next obvious things to be done? What hints can you give to the next person who is going to follow up on your work?

6.4 Reflections

What are the relevant economic, social, environmental, and ethical aspects of your work?

References

- Alif, M. A. R. and Hussain, M. (2025). Yolov12: A breakdown of the key architectural features. *arXiv preprint*, arXiv:2502.14740v1. arXiv.org perpetual non-exclusive license.
- Angulo, A., Vega-Fernández, J. A., Aguilar-Lobo, L. M., Natraj, S., and Ochoa-Ruiz, G. (2019). *Road Damage Detection Acquisition System Based on Deep Neural Networks for Physical Asset Management*, page 3–14. Springer International Publishing.
- Barney, J. (1991). Firm resources and sustained competitive advantage. *Journal of Management*, 17(1):99–120.
- Bochkovskiy, A., Wang, C.-Y., and Liao, H.-Y. M. (2020). Yolov4: Optimal speed and accuracy of object detection. arXiv:2004.10934v1, 23 Apr 2020.
- Civelek, M., Krajčík, V., and Ključnikov, A. (2023). The impacts of dynamic capabilities on smes' digital transformation process: The resource-based view perspective. *Oeconomia Copernicana*, 14(4):1367–1392.
- Creswell, J. W. (2014). *Research Design: Qualitative, Quantitative, and Mixed Methods Approaches*. SAGE Publications, 4th edition.
- Goodfellow, I., Bengio, Y., and Courville, A. (2016). *Deep Learning*. MIT Press.
- Hanelt, A., Bohnsack, R., Marz, D., and Antunes Marante, C. (2020). A systematic review of the literature on digital transformation: Insights

- and implications for strategy and organizational change. *Journal of Management Studies*.
- He, Z., Wang, K., Fang, T., Su, L., Chen, R., and Fei, X. (2024). Comprehensive performance evaluation of yolov11, yolov10, yolov9, yolov8 and yolov5 on object detection of power equipment.
- Hevner, A. R., March, S. T., Park, J., and Ram, S. (2004). Design science in information systems research. *MIS Quarterly*, 28(1):75–105.
- Hidayatullah, P., Syakrani, N., Sholahuddin, M. R., Gelar, T., and Tubagus, R. (2025). Yolov8 to yolo11: A comprehensive architecture in-depth comparative review.
- Jobin, A., Ienca, M., and Vayena, E. (2019). The global landscape of ai ethics guidelines. *Nature Machine Intelligence*, 1:389–399.
- Johnson, R. B. and Onwuegbuzie, A. J. (2004). Mixed methods research: A research paradigm whose time has come. *Educational Researcher*, 33(7):14–26.
- Karbouj, B., Topalian-Rivas, G. A., and Krüger, J. (2024). Comparative performance evaluation of one-stage and two-stage object detectors for screw head detection and classification in disassembly processes. *Procedia CIRP*, 122:527–532. 31st CIRP Conference on Life Cycle Engineering (LCE 2024).
- Khanam, R. and Hussain, M. (2024). Yolov11: An overview of the key architectural enhancements.
- Khanam, R., Hussain, M., Hill, R., and Allen, P. (2024a). A comprehensive review of convolutional neural networks for defect detection in industrial applications. *IEEE Access*, 12:94250–94295.
- Khanam, R., Hussain, M., Hill, R., and Allen, P. (2024b). A comprehensive review of convolutional neural networks for defect detection in industrial applications. *IEEE Access*, 12:94250–94295.
- Krogh, P. (2009). *The DAM book: Digital asset management for photographers*. O’Reilly, Sebastopol, California, 2nd edition.
- Li, C., Li, L., Jiang, H., Weng, K., Geng, Y., Li, L., Ke, Z., Li, Q., Cheng, M., Nie, W., Li, Y., Zhang, B., Liang, Y., Zhou, L., Xu, X., Chu, X., Wei, X., and Wei, X. (2022). Yolov6: A single-stage object detection framework for industrial applications.
- Love, P. E. and Matthews, J. (2019). The ‘how’ of benefits management for digital technology: From engineering to asset management. *Automation in Construction*, 107:102930.
- McCain, E., Mara, N., Van Malssen, K., Carner, D., Reilly, B., Willette, K., Schiefer, S., Askins, J., and Buchanan, S. A. (2021). *Endangered but not too late: The state of digital news preservation*. Donald W. Reynolds Journalism Institute, University of Missouri–Columbia Libraries. OpenAccess. Licensed under CC BY 4.0.
- Prince, S. J. (2023). *Understanding Deep Learning*. The MIT Press.
- Rane, N. (2023). Yolo and faster r-cnn object detection for smart industry 4.0 and industry 5.0: applications, challenges, and opportunities. *SSRN Electronic Journal*.
- Redmon, J., Divvala, S., Girshick, R., and Farhadi, A. (2016). You only look once: Unified, real-time object detection. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 779–788.
- Redmon, J. and Farhadi, A. (2016). Yolo9000: Better, faster, stronger.
- Redmon, J. and Farhadi, A. (2018). Yolov3: An incremental improvement. arXiv:1804.02767v1.
- Sapkota, R., Qureshi, R., Flores-Calero, M., Badgular, C., Nepal, U., Poullose, A., Zeno, P., Vaddevolu, U. B. P., Khan, S., Shoman, M., Yan, H., and Karkee, M. (2025). Yolo11 to its genesis: A decadal and comprehensive review of the you only look once (yolo) series. arXiv, 2406(19407v5).
- Soori, M., Arezoo, B., and Dastres, R. (2023). Artificial intelligence, machine learning and deep learning in advanced robotics, a review. *Cognitive Robotics*, 3:54–70.
- Teng, X., Wu, Z., and Yang, F. (2022). Impact of the digital transformation of small- and medium-sized listed companies on performance: Based on a cost-benefit analysis framework. *Journal of Mathematics*, 2022:1–15.
- Tillväxtverket (2021). Små och medelstora företags digitalisering - vad har betydelse? Technical Report 0366, Tillväxtverket. Accessed: 2025-02-15.
- Ultralytics (2020). Comprehensive guide to ultralytics yolov5. Accessed: 21 February 2025.
- Ultralytics (2023). Yolov8: A unified architecture for object detection, classification, and segmentation. <https://yolov8.com/>. Accessed: 2025-03-01.
- Ultralytics Inc. (2025). Ultralytics YOLO11. <https://docs.ultralytics.com/models/yolo11/>. Accessed: 3 March 2025.

- United Nations (2015). Transforming our world: The 2030 agenda for sustainable development. Accessed: February 28, 2025.
- Verdhan, V. (2021). *Computer Vision Using Deep Learning: Neural Network Architectures with Python and Keras*. Apress, Berkeley, CA, 1st ed. edition.
- Vina, A. (2024). The benefits of ultralytics yolo11 being an anchor-free detector. *Ultralytics Blog*.
- Wang, A., Chen, H., Liu, L., Chen, K., Lin, Z., Han, J., and Ding, G. (2024a). Yolov10: Real-time end-to-end object detection.
- Wang, C.-Y., Bochkovskiy, A., and Liao, H.-Y. M. (2022). Yolov7: Trainable bag-of-freebies sets new state-of-the-art for real-time object detectors. *arXiv preprint arXiv:2207.02696*. Version 1, 6 Jul 2022.
- Wang, C.-Y., Yeh, I.-H., and Liao, H.-Y. M. (2024b). Yolov9: Learning what you want to learn using programmable gradient information. *arXiv preprint arXiv:2402.13616*.
- Wu, M., Brandhorst, H., Marinescu, m.-c., Moré, J., Hlava, M., and Busch, J. (2022). Automated metadata annotation: What is and is not possible with machine learning. *Data Intelligence*, 5:1–17.
- Yin, R. K. (2014). *Case Study Research: Design and Methods*. SAGE Publications, 5th edition.
- Zhang, L., Sun, Z., Tao, H., Wang, M., and Yi, W. (2025). Research on mine-personnel helmet detection based on multi-strategy-improved yolov11. *Sensors (Basel, Switzerland)*, 25(1):170–.

Appendices

A Appendix A: Example Appendix Title

This is an example appendix entry. You can include figures, tables, or additional details relevant to your research.



Figure A-1: An example figure in Appendix A.

Column 1	Column 2
Data 1	Data 2
Data 3	Data 4

Table A.1: An example table in Appendix A.

B Appendix B: Another Appendix Example

You can continue adding appendices in a similar manner.

IEEE Editorial Style Manual: